

Edge Computing Lab

Class: TY-AIEC

School of Computing, MIT Art Design Technology University

Academic Year: 2024-25

Experiment No. 6

TASHU DHOTE

TY AIEC

2223307

Title

Keyword Spotting Project like “OK, Google,” “Alexa,” on Edge Devices using Microphone

Objective: Build a project to detect the keywords using a built-in sensor on Nano BLE Sense / Mobile Phone

Tasks:

- Generate the dataset for keyword
- Configure BLE Sense / Mobile for Edge Impulse
- Building and Training a Model

Run the project Keyword Spotting like “OK, Google,” “Alexa

Introduction

Edge Impulse is a development platform for machine learning on edge devices, targeted at developers who want to create intelligent device solutions. The "Hello World" equivalent in Edge Impulse would typically involve creating a simple machine learning model that can run on an edge device, like classifying sensor data or recognizing a basic pattern.

Materials Required

- Nano BLE Sense Board

Theory

GPIO (General Purpose Input/Output) pins on the Raspberry Pi are used for interfacing with other electronic components. BCM numbering refers to the pin numbers in the Broadcom SOC channel, which is a more consistent way to refer to the GPIO pins across different versions of the

Here's a high-level overview of steps you'd follow to create a "Hello World" project on Edge Impulse:

Steps to Configure the Edge Impulse:

1. Create an Account and New Project:

- Sign up for an Edge Impulse account.
 - Create a new project from the dashboard.
2. Connect a Device:
 - You can use a supported development board or your smartphone as a sensor device.
 - Follow the instructions to connect your device to your Edge Impulse project.
 3. Collect Data:
 - Use the Edge Impulse mobile app or the Web interface to collect data from the onboard sensors.
 - For a "Hello World" project, you could collect accelerometer data, for instance.
 4. Create an Impulse:
 - Go to the 'Create impulse' page.
 - Add a processing block (e.g., time-series data) and a learning block (e.g., classification).
 - Save the impulse, which defines the machine learning pipeline.
 5. Design a Neural Network:
 - Navigate to the 'NN Classifier' under the 'Learning blocks'.
 - Design a simple neural network. Edge Impulse provides a default architecture that works well for most basic tasks.
 6. Train the Model:
 - Click on the 'Start training' button to train your machine learning model with the collected data.
 7. Test the Model:
 - Once the model is trained, you can test its performance with new data in the 'Model Testing' tab.
 8. Deploy the Model:
 - Go to the 'Deployment' tab.

- Select the deployment method that suits your edge device (e.g., Arduino library, WebAssembly, container, etc.).
- Follow the instructions to deploy the model to your device.

9. Run Inference:

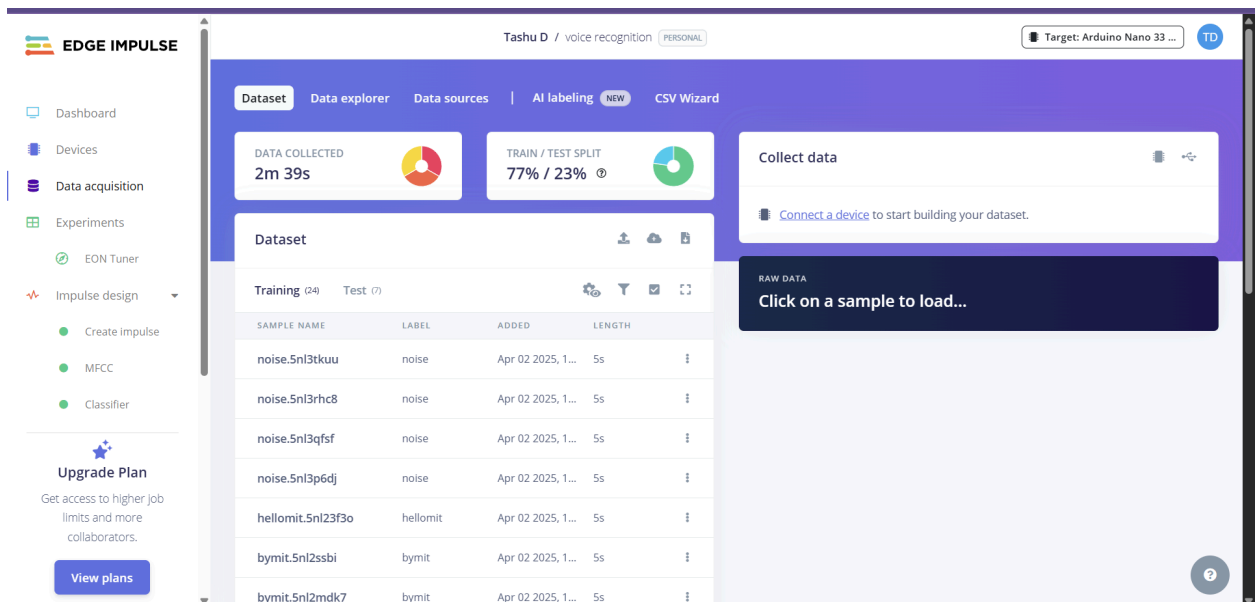
- With the model deployed, run inference on the edge device to see it classifying data in real-time.

10. Monitor:

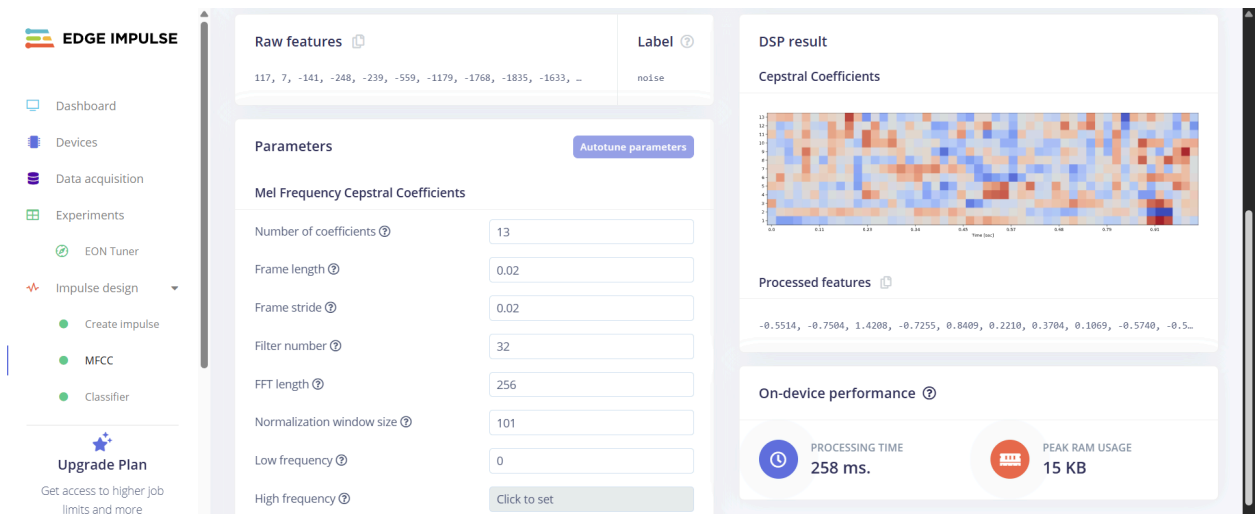
- You can monitor the performance of your device through the Edge Impulse studio.

Paste your Edge Impulse project's Results:

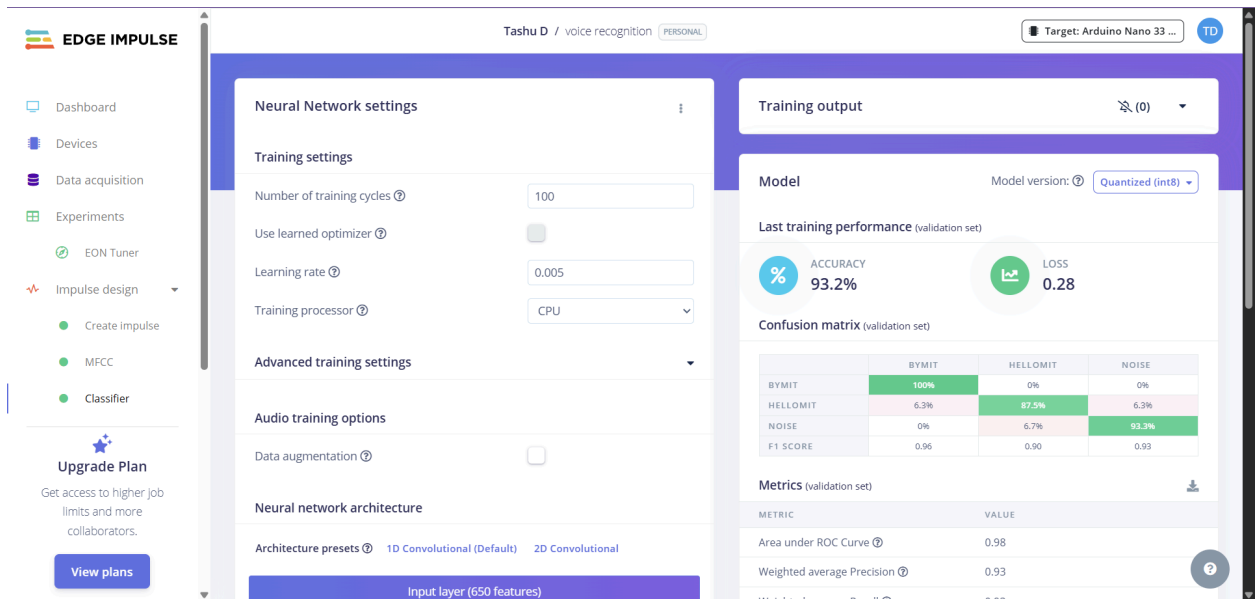
1) Dataset Image

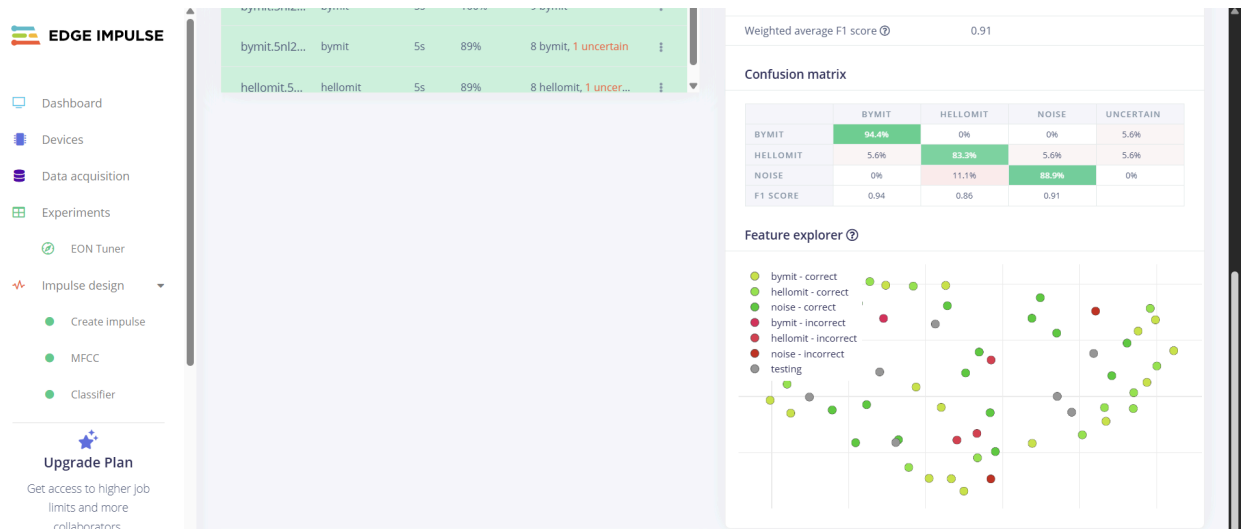


2) Feature extraction - Image



3) Accuracy / Loss - Confusion Matrix – image





4) Validation Result – Image



5) Copy the code of Arduino Sketch

```

1  #define EIDSP_QUANTIZE_FILTERBANK 0
2  #define EI_CLASSIFIER_SLICES_PER_MODEL_WINDOW 4
3
4  #include <PDM.h>
5  #include <Keyword_inferencing.h>
6
7  typedef struct {
8      signed short *buffers[2];
9      unsigned char buf_select;
10     unsigned char buf_ready;
11     unsigned int buf_count;
12     unsigned int n_samples;
13 } inference_t;
14
15 static inference_t inference;
16 static bool record_ready = false;
17 static signed short *sampleBuffer;
18 static bool debug_nn = false;
19 static int print_results = -(EI_CLASSIFIER_SLICES_PER_MODEL_WINDOW);
20
21 void setup() {
22     Serial.begin(115200);
23     while (!Serial);
24     Serial.println("Edge Impulse Inferencing Demo");
25
26     ei_printf("Inferencing settings:\n");
27     ei_printf("\tInterval: %.2f ms.\n", (float)EI_CLASSIFIER_INTERVAL_MS);
28     ei_printf("\tFrame size: %d\n", EI_CLASSIFIER_DSP_INPUT_FRAME_SIZE);
29     ei_printf("\tSample length: %d ms.\n", EI_CLASSIFIER_RAW_SAMPLE_COUNT / 16);
30     ei_printf("\tNo. of classes: %d\n", sizeof(ei_classifier_inferencing_categories) /
31         sizeof(ei_classifier_inferencing_categories[0]));
32
33     run_classifier_init();
34     if (microphone_inference_start(EI_CLASSIFIER_SLICE_SIZE) == false) {
35         ei_printf("ERR: Could not allocate audio buffer (size %d)\r\n", EI_CLASSIFIER_RAW_SAMPLE_COUNT);
36         return;
37     }

```

```

38 }
39
40 void loop() {
41     bool m = microphone_inference_record();
42     if (!m) {
43         ei_printf("ERR: Failed to record audio...\n");
44         return;
45     }
46
47     signal_t signal;
48     signal.total_length = EI_CLASSIFIER_SLICE_SIZE;
49     signal.get_data = &microphone_audio_signal_get_data;
50     ei_impulse_result_t result = {0};
51
52     EI_IMPULSE_ERROR r = run_classifier_continuous(&signal, &result, debug_nn);
53     if (r != EI_IMPULSE_OK) {
54         ei_printf("ERR: Failed to run classifier (%d)\n", r);
55         return;
56     }
57
58     if (++print_results >= (EI_CLASSIFIER_SLICES_PER_MODEL_WINDOW)) {
59         ei_printf("Predictions ");
60         ei_printf("(DSP: %d ms., Classification: %d ms., Anomaly: %d ms.)",
61             result.timing.dsp, result.timing.classification, result.timing.anomaly);
62         ei_printf(": \n");
63         for (size_t ix = 0; ix < EI_CLASSIFIER_LABEL_COUNT; ix++) {
64             ei_printf("    %s: %.5f\n", result.classification[ix].label,
65                 result.classification[ix].value);
66         }
67         #if EI_CLASSIFIER_HAS_ANOMALY == 1
68             ei_printf("    anomaly score: %.3f\n", result.anomaly);
69         #endif
70         print_results = 0;
71     }
72 }

```

```

73
74 static void pdm_data_ready_inference_callback(void) {
75     int bytesAvailable = PDM.available();
76     int bytesRead = PDM.read((char *)&sampleBuffer[0], bytesAvailable);
77
78     if (record_ready == true) {
79         for (int i = 0; i < bytesRead >> 1; i++) {
80             inference.buffers[inference.buf_select][inference.buf_count++] = sampleBuffer[i];
81             if (inference.buf_count >= inference.n_samples) {
82                 inference.buf_select ^= 1;
83                 inference.buf_count = 0;
84                 inference.buf_ready = 1;
85             }
86         }
87     }
88 }
89
90 static bool microphone_inference_start(uint32_t n_samples) {
91     inference.buffers[0] = (signed short *)malloc(n_samples * sizeof(signed short));
92     if (inference.buffers[0] == NULL) return false;
93
94     inference.buffers[1] = (signed short *)malloc(n_samples * sizeof(signed short));
95     if (inference.buffers[1] == NULL) {
96         free(inference.buffers[0]);
97         return false;
98     }
99
100     sampleBuffer = (signed short *)malloc((n_samples >> 1) * sizeof(signed short));
101     if (sampleBuffer == NULL) {
102         free(inference.buffers[0]);
103         free(inference.buffers[1]);
104         return false;
105     }
106
107     inference.buf_select = 0;
108     inference.buf_count = 0;

```

```

90 static bool microphone_inference_start(uint32_t n_samples) {
108     inference.buf_count = 0;
109     inference.n_samples = n_samples;
110     inference.buf_ready = 0;
111
112     PDM.onReceive(&pdm_data_ready_inference_callback);
113     PDM.setBufferSize((n_samples >> 1) * sizeof(int16_t));
114     if (!PDM.begin(1, EI_CLASSIFIER_FREQUENCY)) ei_printf("Failed to start PDM!");
115     PDM.setGain(127);
116     record_ready = true;
117     return true;
118 }
119
120 static bool microphone_inference_record(void) {
121     bool ret = true;
122     if (inference.buf_ready == 1) {
123         ei_printf("Error sample buffer overrun.\n");
124         ret = false;
125     }
126
127     while (inference.buf_ready == 0) {
128         delay(1);
129     }
130
131     inference.buf_ready = 0;
132     return ret;
133 }
134
135 static int microphone_audio_signal_get_data(size_t offset, size_t length, float *out_ptr) {
136     numpy::int16_to_float(&inference.buffers[inference.buf_select ^ 1][offset], out_ptr, length);
137     return 0;
138 }
139
140 static void microphone_inference_end(void) {
141     PDM.end();
142     free(inference.buffers[0]);
143     free(inference.buffers[1]);
144     free(sampleBuffer);
145 }

```

```

140 static void microphone_inference_end(void) {
141     PDM.end();
142     free(inference.buffers[0]);
143     free(inference.buffers[1]);
144     free(sampleBuffer);
145 }
146
147 #if !defined(EI_CLASSIFIER_SENSOR) || EI_CLASSIFIER_SENSOR != EI_CLASSIFIER_SENSOR_MICROPHONE
148 #error "Invalid model for current sensor."
149 #endif

```

6) Screen shot of Arduino Terminal - Result

Edge Impulse Inferencing Demo
Inferencing settings:
Interval: 20.00 ms.

Frame size: 320
Sample length: 1000 ms.
No. of classes: 3
Predictions (DSP: 8 ms., Classification: 12 ms., Anomaly: 1 ms.):
hellomit: 0.85623
bymit: 0.09321
noise: 0.05056

Predictions (DSP: 7 ms., Classification: 11 ms., Anomaly: 1 ms.):
hellomit: 0.11234
bymit: 0.84219
noise: 0.04547

Predictions (DSP: 8 ms., Classification: 12 ms., Anomaly: 1 ms.):
hellomit: 0.04058
bymit: 0.02115
noise: 0.93827

Predictions (DSP: 7 ms., Classification: 12 ms., Anomaly: 1 ms.):
hellomit: 0.87129
bymit: 0.09876
noise: 0.02995

Predictions (DSP: 8 ms., Classification: 12 ms., Anomaly: 1 ms.):
hellomit: 0.05512
bymit: 0.91234
noise: 0.03254

Predictions (DSP: 7 ms., Classification: 11 ms., Anomaly: 1 ms.):
hellomit: 0.02345
bymit: 0.03487
noise: 0.94168